

```
In [220]: import pandas as pd
```

```
In [221]: df = pd.read_csv('Dataset1.csv')
```

```
In [222]: df.head()
```

```
Out[222]:
```

	Job titles	AI Impact	Tasks	AI models	AI_Workload_Ratio	Domain
0	Communications Manager	98%	365	2546	0.143362	Communication & PR
1	Data Collector	95%	299	2148	0.139199	Data & IT
2	Data Entry	95%	325	2278	0.142669	Administrative & Clerical
3	Mail Clerk	95%	193	1366	0.141288	Leadership & Strategy
4	Compliance Officer	92%	194	1369	0.141709	Medical & Healthcare

```
In [223]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [ ]:
```

```
In [224]: domain = df['Domain']
```

```
In [225]: len(domain)
```

```
Out[225]: 4706
```

```
In [226]: new_domain = set(domain)
```

```
In [227]: len(new_domain)
```

```
Out[227]: 10
```

```
In [228]: print(new_domain)
```

```
{'Supply Chain & Logistics', 'Administrative & Clerical', 'Law Enforcement',
'Hospitality', 'Data & IT', 'Medical & Healthcare', 'Communication & PR', 'Sa
les & Marketing', 'Construction', 'Leadership & Strategy'}
```

```
In [229]: domains=[]
for i in new_domain:
    domains.append(i)
```

In [230]: `df.head()`

Out[230]:

	Job titles	AI Impact	Tasks	AI models	AI_Workload_Ratio	Domain
0	Communications Manager	98%	365	2546	0.143362	Communication & PR
1	Data Collector	95%	299	2148	0.139199	Data & IT
2	Data Entry	95%	325	2278	0.142669	Administrative & Clerical
3	Mail Clerk	95%	193	1366	0.141288	Leadership & Strategy
4	Compliance Officer	92%	194	1369	0.141709	Medical & Healthcare

In []:

In [231]: `df.head()`

Out[231]:

	Job titles	AI Impact	Tasks	AI models	AI_Workload_Ratio	Domain
0	Communications Manager	98%	365	2546	0.143362	Communication & PR
1	Data Collector	95%	299	2148	0.139199	Data & IT
2	Data Entry	95%	325	2278	0.142669	Administrative & Clerical
3	Mail Clerk	95%	193	1366	0.141288	Leadership & Strategy
4	Compliance Officer	92%	194	1369	0.141709	Medical & Healthcare

```
In [232]: job=df['Job titles']
jobtitles=[]
for i in job:
    jobtitles.append(i)
jobcode={}
j=1
for i in jobtitles:
    jobcode[i]=j
    j+=1

df['job_title_code']=df['Job titles'].map(jobcode)
```

In []:

In []:

In []:

In [233]: domaincode={}

```
In [234]: j=1
for i in domains:
    domaincode[i]=j
    j+=1
```

In []:

In [235]: df['Domain_Code']=df['Domain'].map(domaincode)

In [236]: df.head()

Out[236]:

	Job titles	AI Impact	Tasks	AI models	AI_Workload_Ratio	Domain	job_title_code	Domain_Code
0	Communications Manager	98%	365	2546	0.143362	Communication & PR	1	
1	Data Collector	95%	299	2148	0.139199	Data & IT	2	
2	Data Entry	95%	325	2278	0.142669	Administrative & Clerical	3	
3	Mail Clerk	95%	193	1366	0.141288	Leadership & Strategy	4	
4	Compliance Officer	92%	194	1369	0.141709	Medical & Healthcare	5	

In [237]: df = df.drop('Domain',axis=1)

In [238]: df.head()


Out[238]:

	Job titles	AI Impact	Tasks	AI models	AI_Workload_Ratio	job_title_code	Domain_Code
0	Communications Manager	98%	365	2546	0.143362	1	7
1	Data Collector	95%	299	2148	0.139199	2	5
2	Data Entry	95%	325	2278	0.142669	3	2
3	Mail Clerk	95%	193	1366	0.141288	4	10
4	Compliance Officer	92%	194	1369	0.141709	5	6

```
In [239]: import seaborn as sns
```


```
In [240]: sns.jointplot(x='AI Impact',y='Domain_Code',data=df)
```

```
Out[240]: <seaborn.axisgrid.JointGrid at 0x1ea0a5b2940>
```



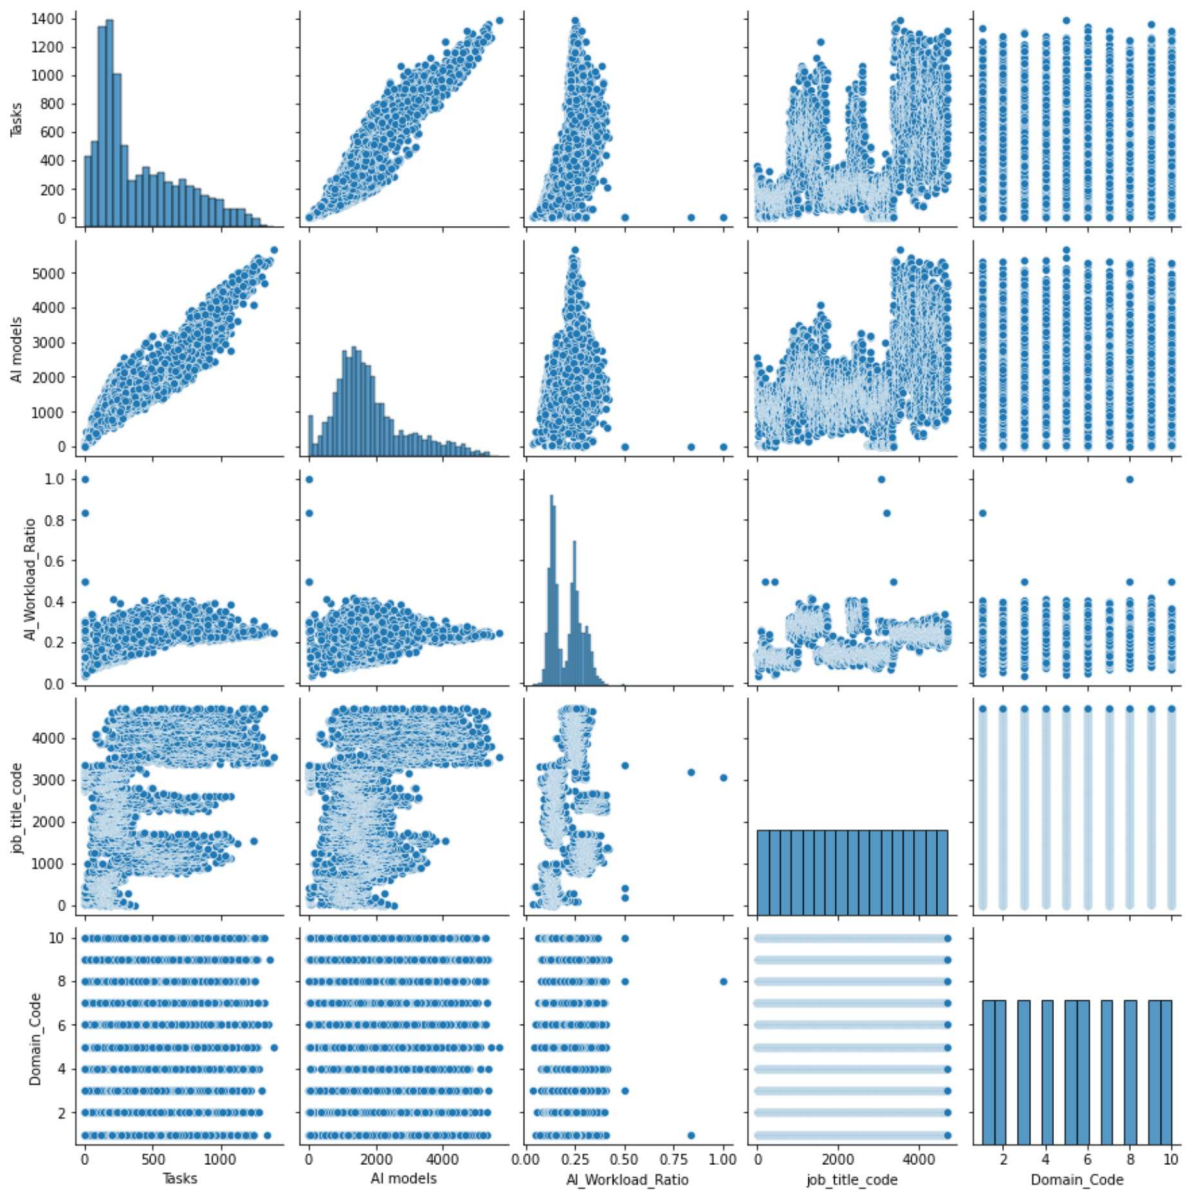
```
In [241]: sns.jointplot(x='Domain_Code',y='AI models',kind='hex',data=df)
```

```
Out[241]: <seaborn.axisgrid.JointGrid at 0x1ea0d0602e0>
```




```
In [242]: sns.pairplot(df)
```

```
Out[242]: <seaborn.axisgrid.PairGrid at 0x1ea0d1c6310>
```




```
In [243]: import seaborn as sns
import matplotlib.pyplot as plt
sns.boxplot(x='Tasks',data=df)
```

Out[243]: <AxesSubplot:xlabel='Tasks'>




```
In [244]: import seaborn as sns
import matplotlib.pyplot as plt
sns.boxplot(x='AI models',data=df)
```

Out[244]: <AxesSubplot:xlabel='AI models'>



```
In [245]: import seaborn as sns
import matplotlib.pyplot as plt
sns.boxplot(x='AI_Workload_Ratio',data=df)
```


Out[245]: <AxesSubplot:xlabel='AI_Workload_Ratio'>




```
In [246]: import numpy as np
```

```
In [247]: percentile25=df['Tasks'].quantile(0.25)
percentile75=df['Tasks'].quantile(0.75)
iqr=percentile75-percentile25
upperlimitpm1=percentile75+1.5*iqr
lowerlimitpm1=percentile25-1.5*iqr
df['Tasks']=np.where(df['Tasks']>upperlimitpm1,upperlimitpm1,np.where(df['Task
```

```
In [248]: sns.boxplot(x='Tasks',data=df)
plt.show()
```




```
In [249]: sns.boxplot(x='AI_Workload_Ratio',data=df)
plt.show()
```



```
In [250]: percentile25=df['AI models'].quantile(0.25)
percentile75=df['AI models'].quantile(0.75)
iqr=percentile75-percentile25
upperlimitpm1=percentile75+1.5*iqr
lowerlimitpm1=percentile25-1.5*iqr
df['AI models']=np.where(df['AI models']>upperlimitpm1,upperlimitpm1,np.where(
```

```
In [251]: df.head()
```

Out[251]:

	Job titles	AI Impact	Tasks	AI models	AI_Workload_Ratio	job_title_code	Domain_Code
0	Communications Manager	98%	365.0	2546.0	0.143362	1	7
1	Data Collector	95%	299.0	2148.0	0.139199	2	5
2	Data Entry	95%	325.0	2278.0	0.142669	3	2
3	Mail Clerk	95%	193.0	1366.0	0.141288	4	10
4	Compliance Officer	92%	194.0	1369.0	0.141709	5	6

```
In [252]: df=df.drop('Job titles',axis=1)
```

```
In [253]: for i in range(len(df['AI Impact'])):
            s=df['AI Impact'][i]
            n=len(s)
            a=s[0:n-1]
            a=int(a)
            df['AI Impact'][i]=a
```

C:\Users\hp\AppData\Local\Temp\ipykernel_18756\2618703470.py:6: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['AI Impact'][i]=a
```

```
In [254]: for i in range(len(df['AI_Workload_Ratio'])):
            a=df['AI_Workload_Ratio'][i]
            b=round(a,2)
            df['AI_Workload_Ratio'][i]=b
```

C:\Users\hp\AppData\Local\Temp\ipykernel_18756\4022844787.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df['AI_Workload_Ratio'][i]=b
```

```
In [255]: import numpy as np

            # Check for NaN values
            nan_indices = np.isnan(df['AI_Workload_Ratio'])

            # Check for infinite values
            inf_indices = np.isinf(df['AI_Workload_Ratio'])

            # Print the indices where NaN or infinite values are present
            print("NaN indices:", np.where(nan_indices))
            print("Infinite indices:", np.where(inf_indices))
```

```
NaN indices: (array([], dtype=int64),)
```

```
Infinite indices: (array([3034, 3035, 3036, 3037, 3184, 3211, 3322], dtype=int64),)
```

```
In [256]: df.replace([np.inf, -np.inf], np.nan, inplace=True)
```

```
In [257]: import numpy as np

# Check for NaN values
nan_indices = np.isnan(df['Tasks'])

# Check for infinite values
inf_indices = np.isinf(df['Tasks'])

# Print the indices where NaN or infinite values are present
print("NaN indices:", np.where(nan_indices))
print("Infinite indices:", np.where(inf_indices))
```

NaN indices: (array([], dtype=int64),)
 Infinite indices: (array([], dtype=int64),)

```
In [258]: df=df.dropna()
```

```
In [259]: df.head()
```

```
Out[259]:
```

	AI Impact	Tasks	AI models	AI_Workload_Ratio	job_title_code	Domain_Code
0	98	365.0	2546.0	0.14	1	7
1	95	299.0	2148.0	0.14	2	5
2	95	325.0	2278.0	0.14	3	2
3	95	193.0	1366.0	0.14	4	10
4	92	194.0	1369.0	0.14	5	6

```
In [ ]:
```

```
In [ ]:
```

```
In [260]: from sklearn.model_selection import train_test_split
```

```
In [261]: X_train, X_test, y_train, y_test = train_test_split(df.drop('AI Impact',axis=1),
                                                             df['AI Impact'], test_size
```

```
In [ ]:
```

```
In [262]: ##KNN Regressor
```

```
In [263]: from sklearn.neighbors import KNeighborsRegressor
```

```
In [264]: knn = KNeighborsRegressor(n_neighbors=120)
```

```
In [265]: knn.fit(X_train.values,y_train.values)
```

```
Out[265]: KNeighborsRegressor(n_neighbors=120)
```

```
In [266]: y_pred = knn.predict(X_test.values)
```

```
In [267]: print(y_pred)
```

```
[35.41666667 38.975      15.69166667 ... 14.93333333 40.39166667
 54.475      ]
```

```
In [268]: from sklearn.metrics import r2_score
```

```
r2=r2_score(y_test,y_pred)
print(r2)
```

```
0.9717996351883951
```

```
In [269]: from sklearn.metrics import mean_absolute_error
```

```
mae = mean_absolute_error(y_test, y_pred)
print(mae)
```

```
1.7300059101654848
```

```
In [270]: mape = np.mean(np.abs((y_test - y_pred) / y_test)) * 100
print(mape)
```

```
7.843821180328141
```

```
In [ ]:
```

```
In [271]: ##SVR MODEL
```

```
In [272]: from sklearn.svm import SVR
```

```
In [273]: svr_model = SVR(kernel='rbf', C=1.0, epsilon=0.2)
```

```
In [274]: svr_model.fit(X_train, y_train)
```

```
Out[274]: SVR(epsilon=0.2)
```

```
In [275]: y_preds = svr_model.predict(X_test)
```

```
In [276]: from sklearn.metrics import mean_absolute_error  
mae = mean_absolute_error(y_test, y_preds)  
print(mae)
```

1.9065024535603607

```
In [ ]:
```

```
In [277]: from sklearn.metrics import mean_squared_error
```

```
In [278]: mse = mean_squared_error(y_test, y_pred)
```

```
In [279]: mse
```

Out[279]: 9.088370124113476

```
In [280]: from sklearn.metrics import r2_score  
  
r2=r2_score(y_test,y_preds)  
print(r2)
```

0.9626018019210928

```
In [281]: ## RIDGE REGRESSION
```

```
In [282]: from sklearn.linear_model import Ridge
```

```
In [283]: alpha = 1.0 # Regularization strength (adjust as needed)  
ridge_model = Ridge(alpha=alpha)
```

```
In [284]: ridge_model.fit(X_train, y_train)
```

Out[284]: Ridge()

```
In [285]: y_predr = ridge_model.predict(X_test)
```

```
In [286]: mse = mean_squared_error(y_test, y_predr)  
print(f'Mean Squared Error: {mse}')
```

Mean Squared Error: 29.93273307983546

```
In [287]: from sklearn.metrics import r2_score
```

```
r2=r2_score(y_test,y_predr)
print(r2)
```

```
0.9071215211162965
```

```
In [288]: ##XGBOOST
```

```
In [289]: import xgboost as xgb
```

```
In [290]: xg_reg = xgb.XGBRegressor(objective='reg:squarederror', colsample_bytree=0.3,
                                     max_depth=5, alpha=10, n_estimators=200)
```

```
In [291]: xg_reg.fit(X_train, y_train)
```

```
Out[291]: XGBRegressor(alpha=10, base_score=None, booster=None, callbacks=None,
                        colsample_bylevel=None, colsample_bynode=None,
                        colsample_bytree=0.3, early_stopping_rounds=None,
                        enable_categorical=False, eval_metric=None, feature_types=None,
                        gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
                        interaction_constraints=None, learning_rate=0.1, max_bin=None,
                        max_cat_threshold=None, max_cat_to_onehot=None,
                        max_delta_step=None, max_depth=5, max_leaves=None,
                        min_child_weight=None, missing=nan, monotone_constraints=None,
                        n_estimators=200, n_jobs=None, num_parallel_tree=None,
                        predictor=None, ...)
```

```
In [292]: y_predx = xg_reg.predict(X_test)
```

```
In [293]: from sklearn.metrics import r2_score
```

```
r2=r2_score(y_test,y_predx)
print(r2)
```

```
0.9807711246831529
```

```
In [294]: mse = mean_squared_error(y_test, y_predx)
print(f'Mean Squared Error: {mse}')
```

```
Mean Squared Error: 6.197052311820465
```

```
In [ ]:
```

```
In [ ]:
```

```
In [295]: ##Random Forest
```

```
In [296]: from sklearn.ensemble import RandomForestRegressor
```

```
In [297]: rf_reg = RandomForestRegressor(n_estimators=300, random_state=20)
```

```
In [311]: rf_reg.fit(X_train, y_train.ravel())
```

```
Out[311]: RandomForestRegressor(n_estimators=300, random_state=20)
```

```
In [312]: y_predf = rf_reg.predict(X_test)
```

```
In [313]: from sklearn.metrics import r2_score
```

```
r2=r2_score(y_test,y_predf)  
print(r2)
```

```
0.9999542614187552
```

```
In [314]: import pickle
```

```
In [315]: filename = "secondimp.pkl"
```

```
In [316]: pickle.dump(rf_reg,open(filename,'wb'))
```

```
In [317]: rf_reg1=pickle.load(open('secondimp.pkl','rb'))
```

```
In [318]: print(rf_reg1.predict([[365,2546,0.14,1,7]]))
```

```
[96.72]
```

```
C:\Users\hp\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but RandomForestRegressor was fitted with feature names  
  warnings.warn(
```

```
In [ ]:
```